

SIMATS ENGINEERING

ASSESSMENT TOOL 2

COURSE NAME: Artificial Intelligence

COURSE CODE: CSA1711

COURSE OUTCOME COVERED

CO4: *Implement machine learning techniques including decision tree algorithms, reinforcement learning, and build models for classification and decision-making. (BL3)*

Assessment Tool Weightage: 50%

QUESTIONS: 4

Duration: 1 Hour

Total: Marks:40

Annexure A

5

Code of Conduct:

I Kishore S.P (192511081) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment.

I understand that any violation of academic integrity rules will result in disciplinary action.

Signature of the student:

INDUSTRY PROBLEM TASK

INDUSTRY SCENARIO

Context: A healthcare company has collected patient records (age, blood pressure, cholesterol, glucose level, BMI, family history) to build a predictive model for early diagnosis of diabetes. The company also wants to train an AI agent to recommend personalised treatment actions (Diet / Exercise / Medication / Monitor) based on patient state transitions over time.

You are hired as an AI/ML Engineer. Address the following tasks:

Task 1: Data Preparation & Inductive Learning

- (a) Describe the inductive learning approach suitable for this medical dataset. Identify the target variable and at least three key features with justification.
- (b) Explain how you would handle class imbalance (more non-diabetic than diabetic cases) in the dataset. Suggest and justify a suitable balancing strategy (SMOTE / under-sampling / class weights).
- (c) Split the dataset (70:30) and justify the split ratio for a medical use-case. State the preprocessing steps (normalisation, encoding) applied and their impact.

Task 2: Decision Tree for Diagnosis

- (a) Build a Decision Tree classifier and draw the first three levels of the tree. Justify the root node selection using Information Gain / Gini Index values.
- (b) Evaluate the model: report Accuracy, Precision, Recall, and F1-Score. Identify False Negatives (missed diabetes cases) and suggest how to reduce them—justify clinically.
- (c) Apply post-pruning (cost-complexity pruning). Compare pre-pruning vs post-pruning accuracy and explain which model is more appropriate for deployment in a clinical setting.

Task 3: Statistical Learning for Risk Stratification

- (a) Apply Naïve Bayes or Logistic Regression as a statistical learning model on the same dataset. Compare its performance (Accuracy, AUC-ROC) with the Decision Tree. Discuss when a statistical model is preferable.
- (b) Perform feature importance analysis. Identify the top 3 predictors of diabetes from both the Decision Tree and the statistical model. Do they agree? Justify your findings.

Task 4: Reinforcement Learning for Treatment Recommendation

- (a) Model the treatment recommendation as an MDP. Define: States (patient health stages), Actions (Diet / Exercise / Medication / Monitor), Reward function (health improvement score), and Transition dynamics. Justify each component.
- (b) Apply Q-Learning to train the agent for at least 5 patient episodes. Show the Q-table update for at least 3 state-action pairs across 2 iterations. State the convergence criterion used.
- (c) Extract the final learned policy. Discuss ethical considerations of deploying a Reinforcement Learning agent for medical treatment recommendations (patient safety, bias, explainability).

Task 1: Data Preparation & Inductive Learning

(a) Inductive Learning

Use supervised inductive learning to learn the relationship between patient features and diabetes status.

- Target: Diabetes (Yes/No)
- Key features:
 - Glucose: strong indicator of diabetes.
 - BMI: associated with insulin resistance.
 - Age: diabetes risk generally increases with age.
 - Blood pressure/family history: useful additional risk factors.

(b) Class Imbalance

Use SMOTE on the training set because diabetic cases are fewer. It creates synthetic minority examples and improves recall without discarding many non-diabetic records.

(c) 70:30 Split & Preprocessing

Use a 70% training / 30% testing stratified split. This provides enough data for training while keeping a sufficiently large test set.

Preprocessing:


- Encode Family History: Yes = 1, No = 0.
- Handle missing values using imputation.
- Standardise numerical features for Logistic Regression.
- Decision Trees do not require normalisation.
- Apply SMOTE only to training data.

Task 2: Decision Tree

(a) Tree Construction

Use a Decision Tree with Gini Index or Information Gain.

Example:



The feature with the highest Information Gain / lowest Gini impurity becomes the root. Typically, glucose is expected to be a strong root feature, but the actual dataset must determine the result.

(b) Evaluation

Use:

- Accuracy = correct predictions / total predictions
- Precision = $TP / (TP + FP)$
- Recall = $TP / (TP + FN)$
- F1-score = harmonic mean of precision and recall.

False Negative: diabetic patient predicted as non-diabetic.

To reduce false negatives:

- use class weights;
- apply SMOTE;
- lower the classification threshold where clinically appropriate;
- prioritise high recall/sensitivity.

This is important because missed diabetes cases can delay diagnosis and treatment.

(c) Post-Pruning

Use cost-complexity pruning:

$$R_{\alpha}(T) = R(T) + \alpha |T|$$

It removes unnecessary branches and reduces overfitting.

A pruned tree is generally preferable for clinical deployment if it maintains or improves validation performance while being simpler and more interpretable.

Task 3: Statistical Learning

(a) Logistic Regression

Use Logistic Regression as a statistical baseline.

$$P(\text{Diabetes}) = \frac{1}{1 + e^{-z}}$$

Compare it with the Decision Tree using:

- Accuracy
- AUC-ROC

A statistical model is preferable when interpretability, probability estimates, and understanding predictor effects are important.

(b) Feature Importance

Possible top predictors:

Rank	Decision Tree	Logistic Regression
1	Glucose	Glucose
2	BMI	BMI
3	Age	Age

If both models agree, this suggests these variables contain strong predictive information. However, feature importance does not prove causation.

Task 4: Reinforcement Learning

(a) MDP

$$MDP = (S, A, P, R, \gamma)$$

- States: Low risk, Moderate risk, High risk, Poor control, Improved.
- Actions: Diet, Exercise, Medication review, Monitor.
- Reward: Positive for health improvement and negative for deterioration/adverse outcomes.
- Transitions: Patient moves between health states according to treatment and health changes.

(b) Q-Learning

Update rule:

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

Example with ($\alpha=0.5, \gamma=0.9$):

If:

$$Q(S_2, D) = 0, \quad r = 5$$

then:

$$[Q(S_2, D) = 0 + 0.5(5) = 2.5]$$

After another update where the next state's best Q-value is 2.5:

$$[Q(S_2, D) = 2.5 + 0.5[5 + 0.9(2.5) - 2.5] = 4.875]$$

Another example:

$$[Q(S_3, \text{Monitor}) = 0.5[2 + 0.9(4.875)] = 3.194]$$

Train for at least 5 patient episodes.

Convergence criterion:

$$[\max |Q_{\text{new}} - Q_{\text{old}}| < 0.001]$$

for several consecutive iterations.

(c) Final Policy

The policy is:

$$[\pi(s) = \arg \max_a Q(s, a)]$$

Example:

Patient State	Best Action
Low risk	Exercise
Moderate risk	Exercise
High risk	Diet
Poor control	Medication review
Improved	Monitor

Ethical Considerations

- Safety: RL must not prescribe medication autonomously.
- Bias: evaluate performance across patient groups.
- Explainability: provide reasons for recommendations.
- Human oversight: clinicians must approve treatment decisions.
- Privacy: protect patient medical data.
- Validation: test extensively before clinical deployment.

Conclusion

A suitable solution is Decision Tree + Logistic Regression for diabetes prediction, followed by a constrained Q-Learning framework for treatment recommendations. In clinical deployment, recall, safety, fairness and interpretability are more important than accuracy alone.

RUBRICS FOR CALCULATION:

Criteria	Weightage	Level 4 – Exemplary	Level 3 – Proficient	Level 2 – Developing	Level 1 – Beginning
Criteria	Weightage	Level 4 – Exemplary	Level 3 – Proficient	Level 2 – Developing	Level 1 – Beginning
Understanding of Industry Problem	20%	Comprehensive understanding of real-world problem and constraints	Good understanding with minor gaps in requirements	Basic understanding; some requirements unclear	Poor understanding; misinterprets problem context
Application of Engineering Concepts	25%	Applies advanced and relevant concepts effectively to solve the problem	Applies appropriate concepts with minor errors	Limited application of concepts	Incorrect or no application of concepts
Solution Design	25%	Innovative, practical, and feasible solution aligned with industry standards	Logical and feasible solution with minor gaps	Basic solution with limited feasibility	Unclear or impractical solution
Use of Tools	15%	Effective use of modern tools, technologies, and real data	Appropriate use of tools with correct implementation	Limited or inconsistent use of tools	Minimal or incorrect use of tools
Analysis	10%	Thorough analysis with validated results and performance evaluation	Good analysis with reasonable validation	Basic analysis with limited validation	Poor or no analysis
Professionalism	5%	Highly professional, well-documented, and clearly presented work	Good documentation and presentation	Basic documentation with some gaps	Poor documentation and presentation